

Bharat-AI-SoC-Student-Challenge

Team Members

- Hari Kesavaraj J
- Iniyavan S
- Karthiswaran R

3rd Year B.E Electronics Engineering
(VLSI Design & Technology)
K. S. Rangasamy College of Technology

GitHub: <https://github.com/Karthiswaran-R/Bharat-AI-SoC-Student-Challenge>

Project Focus: Object detection on Xilinx PYNQ-Z2

1. Abstract

This project presents a CPU-only object detection system implemented on the Xilinx PYNQ-Z2, developed under Problem Statement 5 of the Bharat AI-SoC Student Challenge.

The system performs pedestrian detection using a classical computer vision pipeline based on:

- Histogram of Oriented Gradients (HOG)
- Support Vector Machine (SVM)
- Non-Maximum Suppression (NMS)

All computations are executed on the ARM Cortex-A9 processor of the Zynq-7020 SoC without FPGA acceleration.

This implementation establishes a quantitative baseline for future hardware acceleration using the programmable logic (PL) fabric.

2. Introduction

Embedded object detection is critical in:

- Smart surveillance
- Automotive systems
- Edge AI applications
- Smart city infrastructure


However, real-time object detection on resource-constrained SoCs is challenging due to:



- Limited processing power
- Memory bandwidth constraints
- Power limitations

This project evaluates the performance limitations of a CPU-only implementation before introducing FPGA acceleration.

3. System Architecture

3.1 Hardware Platform





Hardware Specifications

Component	Description
SoC	Zynq-7020
Processor	Dual-core ARM Cortex-A9
FPGA	Artix-7 Programmable Logic
RAM	512 MB DDR3
OS	PYNQ Linux

4. Methodology

4.1 Image Preprocessing

- Resize image to fixed width (640 px)
- Preserve aspect ratio
- Improve detection consistency

4.2 Feature Extraction – HOG

HOG computes gradient orientation histograms over localized regions.

Advantages:

- Lightweight
- Deterministic runtime
- Suitable for embedded CPU baseline

Limitations:

- Computationally expensive sliding window
- Not real-time on Cortex-A9

4.3 Classification – SVM

Uses OpenCV's pre-trained pedestrian detector.

Binary classification:

- Person
- Background

5. Experimental Setup

- Images stored in /home/xilinx/jupyter_notebooks/images
- Python 3 + OpenCV
- Detection timing measured using time.time()
- Visualization excluded from latency measurement

6. Performance Evaluation

Metrics Used

- Average Latency (ms)
- Minimum Latency (ms)
- Maximum Latency (ms)
- Frames Per Second (FPS)
- Total Execution Time

Output:



Figure 1



Figure 2

```
===== CPU-ONLY PERFORMANCE (PYNQ-Z2) =====
Total images processed : 8
Average latency        : 8437.73 ms
Min latency            : 3854.21 ms
Max latency            : 10118.12 ms
Throughput (FPS)       : 0.12
Total execution time   : 79.71 s
```

Figure 3

Observed Performance

Metric	Observed Value
Average Latency	~ 4000 – 5000 ms
FPS	~ 0.2
Real-Time Capability	Not Achieved

7. Analysis

7.1 Why Performance is Low

- Sliding window scanning
- Image pyramid scaling

- CPU-only execution
- No SIMD optimization
- No hardware acceleration

The Cortex-A9 processor lacks sufficient computational throughput for real-time object detection using HOG.

8. Importance of Baseline

This CPU-only implementation provides:

- Reference latency
- Reference FPS
- Engineering comparison point

9. Future Work – FPGA Acceleration Plan

Planned hardware offloading:

- Gradient computation module (PL)
- Sliding window engine (PL)
- Feature histogram computation (PL)
- AXI-based PS-PL communication

Target:

- $\geq 5\times$ performance improvement
- ≥ 2 FPS real-time threshold
- Optimized memory transfers

10. Conclusion

This work successfully:

- Implements CPU-only object detection on PYNQ-Z2
- Benchmarks real embedded performance
- Establishes quantitative baseline

- Prepares for hardware acceleration phase

The results confirm that CPU-only detection is insufficient for real-time embedded applications, justifying FPGA-based acceleration.

12. References

- [1] Xilinx Inc., “PYNQ: Python Productivity for Zynq,” *PYNQ Documentation*, 2024. [Online]. Available: <https://pynq.readthedocs.io/en/latest/>.
- [2] OpenCV Team, “Open Source Computer Vision Library (OpenCV) – Version 4.x Documentation,” 2024. [Online]. Available: <https://docs.opencv.org/4.x/>.
- [3] Xilinx Inc., *Zynq-7000 SoC Technical Reference Manual*, UG585, 2023.
- [4] Arm Ltd., *ARM Cortex-A9 Technical Reference Manual*, Arm Documentation, 2023.